

Reinforcement Learning for Autonomous Quadrotor Helicopter Control

Michael Koval, Christopher Mansley, and Michael Littman

Rutgers, the State University of New Jersey

Abstract

- Quadrotor helicopters are rapidly finding use for a number of applications
- Programming a new behavior means implementing a control algorithm
- Tuning a new algorithms is costly and requires expertise in control theory
- Supervised learning is not a viable alternative: training data is hard to get
- Reinforcement learning can learn directly from the quadrotor's raw sensors
- **Goal:** Use reinforcement learning to teach a quadrotor helicopter how to perform complex tasks with minimal human interaction

Experimental Setup

1. Parrot AR.Drone [4]

- Inexpensive quadrotor helicopter with impressive technical specifications
- Outfitted with two cameras, an altitude sensor, and an IMU
- On-board computer running a variant of embedded Linux
- Communicates with a controlling computer over a dedicated wifi network

2. Vicon Motion Tracking System

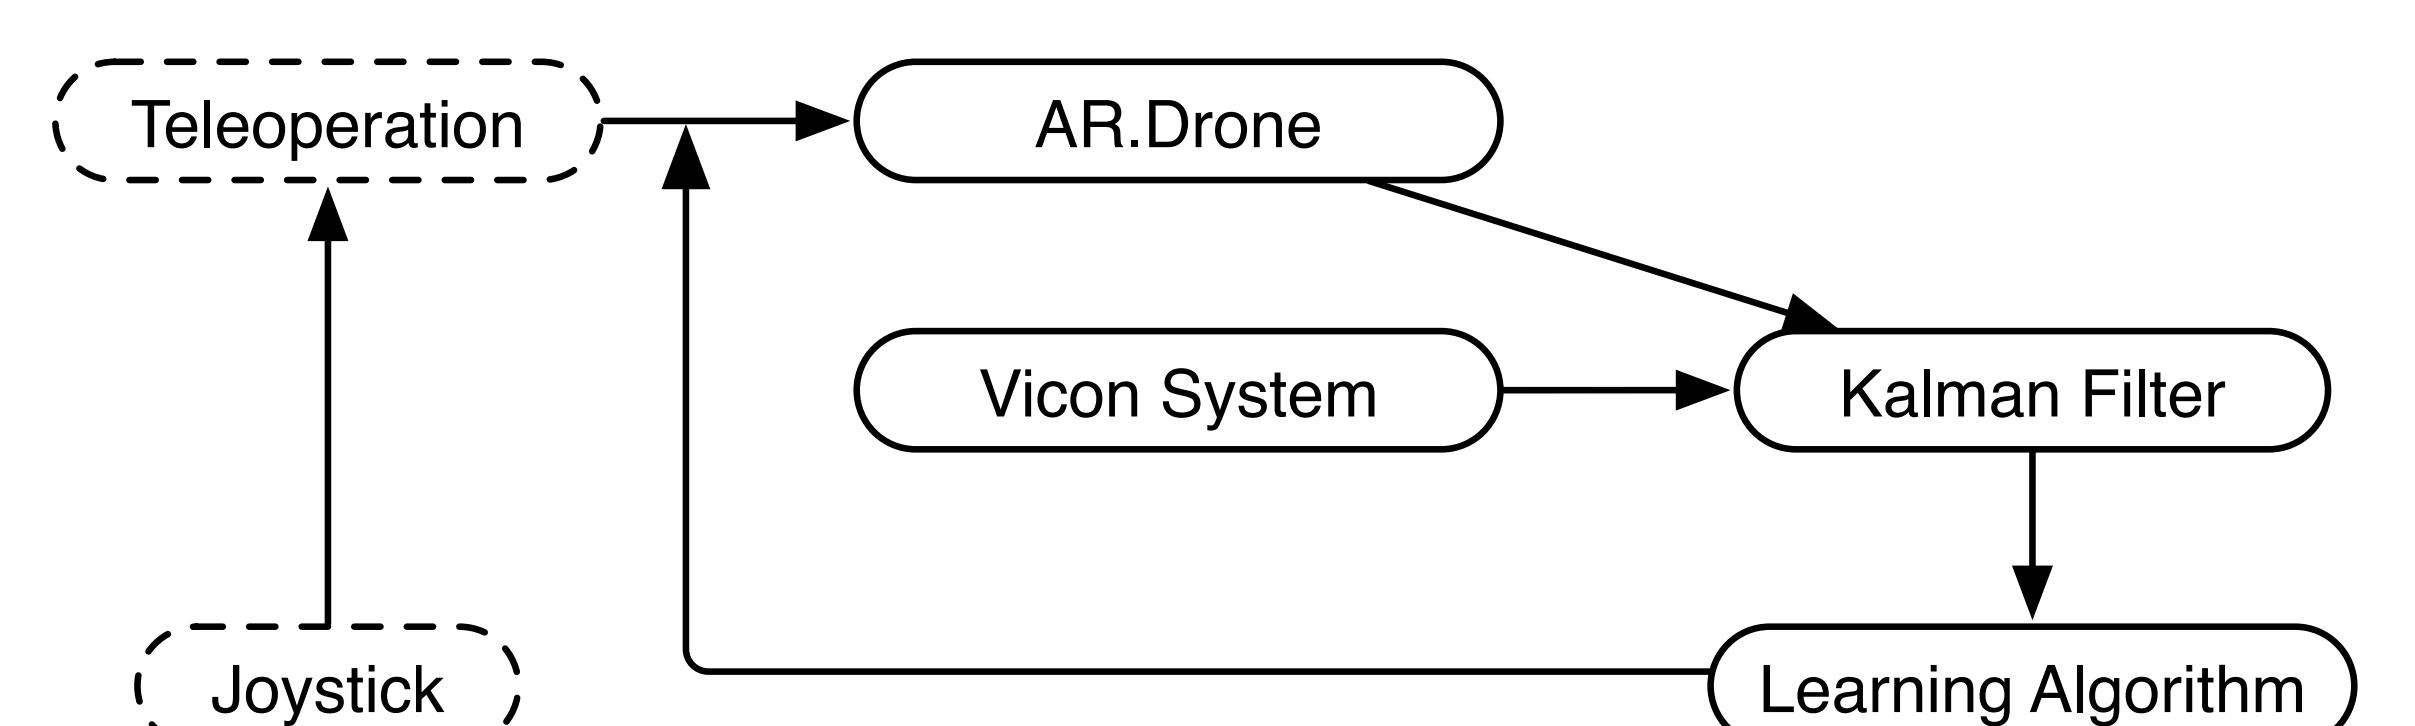
- Quadrotor is tagged with an asymmetric pattern of reflective spheres
- Configure the Vicon system to track the pattern as a rigid body
- Track the quadrotor's flight using the Vicon system's infrared cameras

3. Control Software

- Exerts direct, low-level control over the AR.Drone's flight
- Closes the loop between the motion tracking system and the AR.Drone
- Uses reinforcement learning to find an optimal control algorithm

4. Robot Operating System (ROS)

- Inter-process communication framework developed by Willow Garage
- Allows the system to be easily extended for more complex experiments



Reinforcement Learning

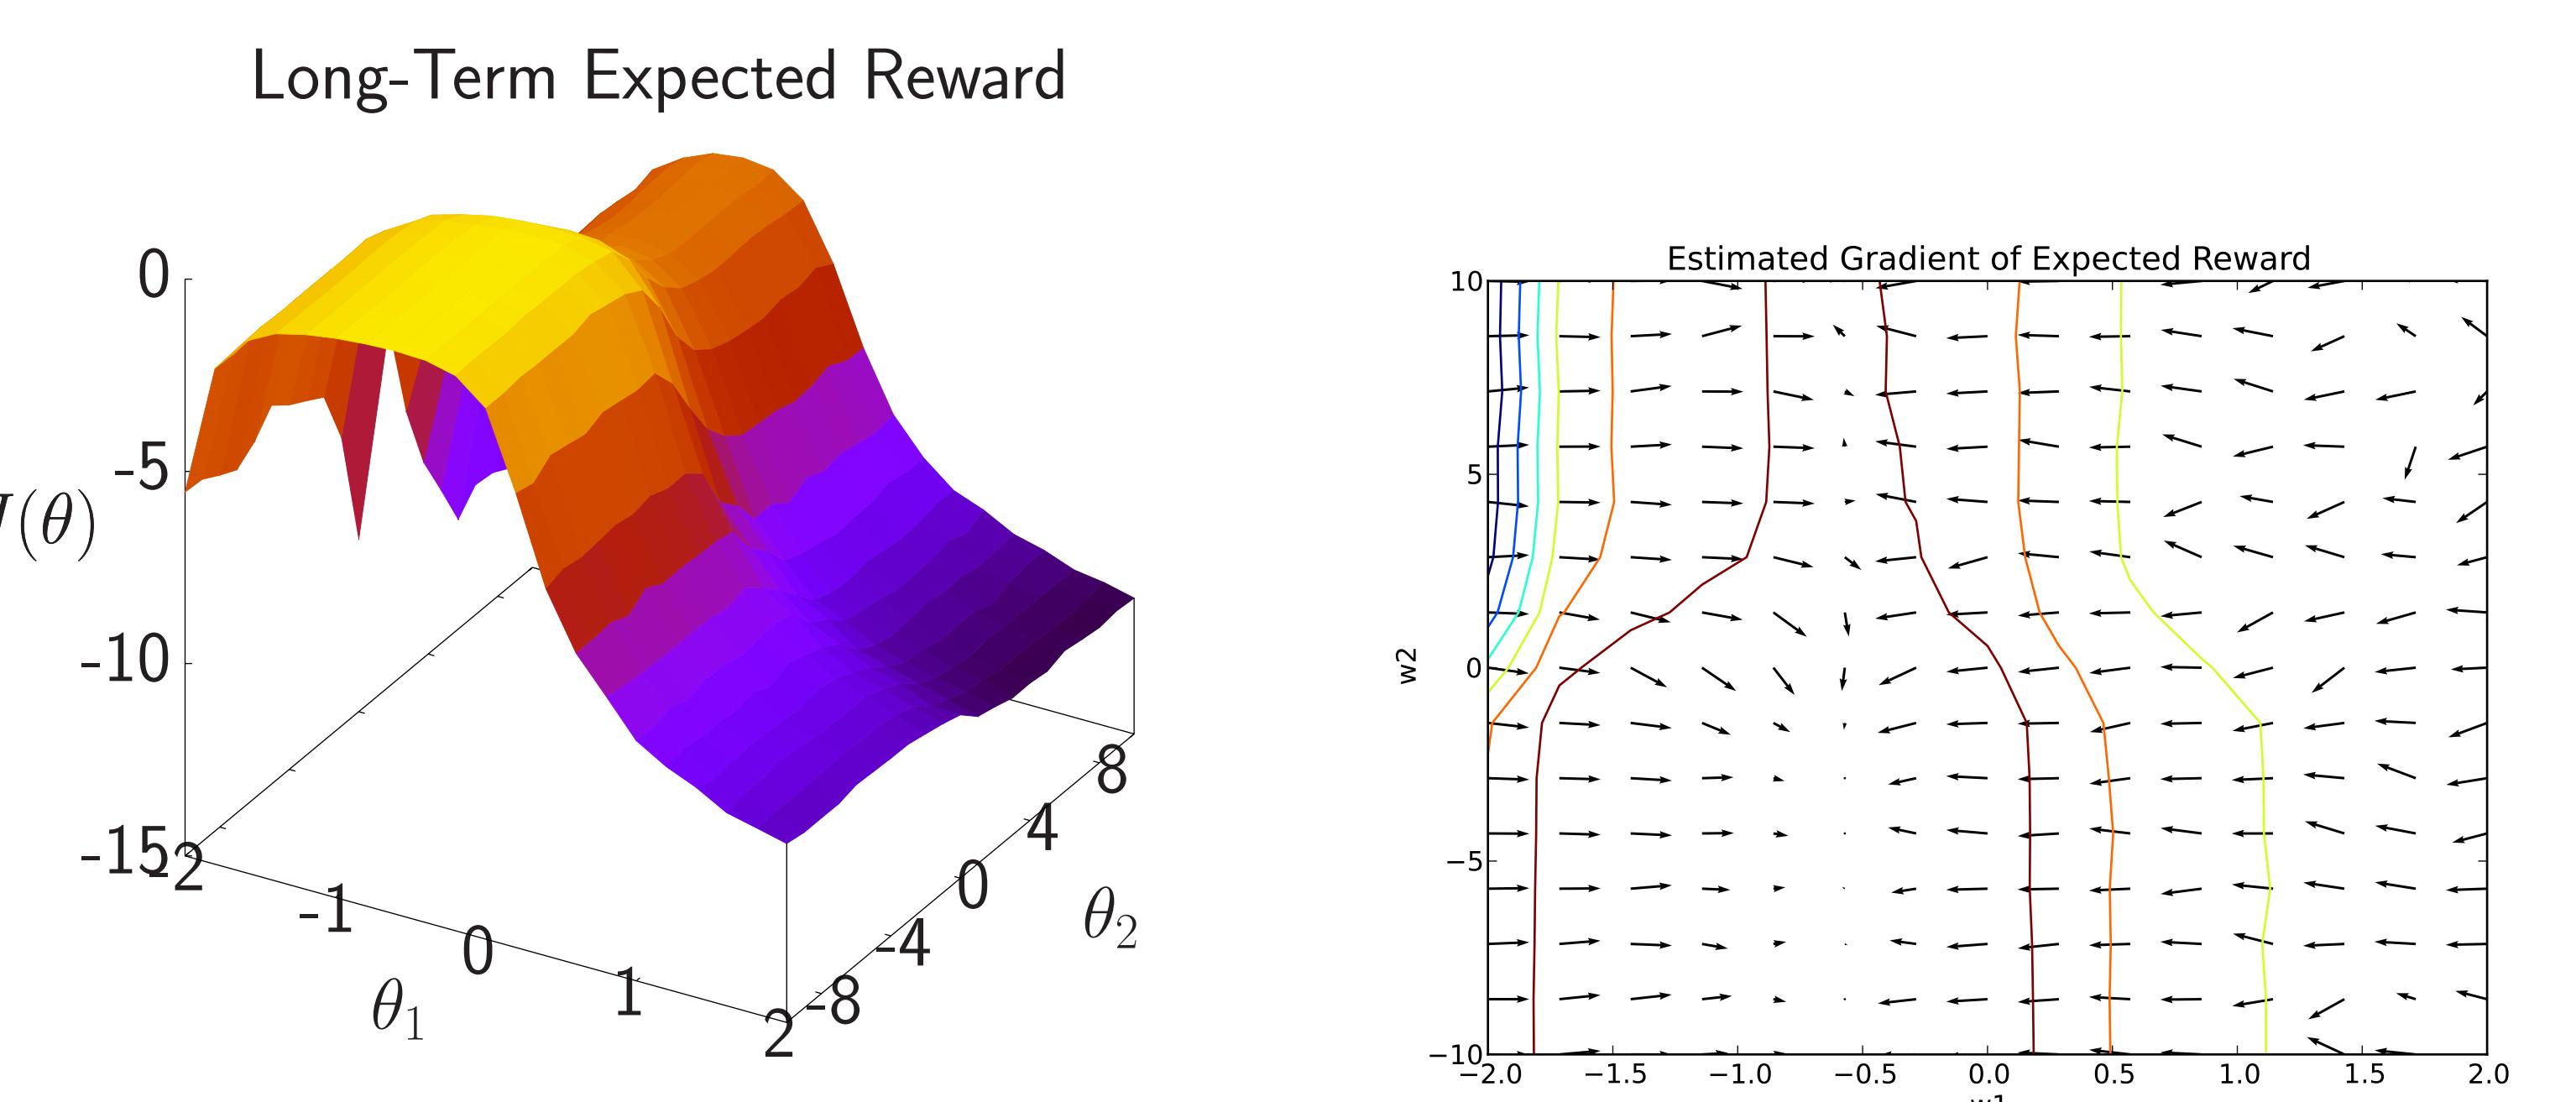
- Agent learns from rewards earned while interacting with a stochastic environment
- Does not require annotated examples of the agent's correct actions
- Ideal for situations rewards can automatically be assigned
- Environment is modeled as a Markov decision process where, at time t :
 1. the agent is in state $s_t \in S$,
 2. chooses action $a_t \sim \pi(s_t, a)$ using policy π , and
 3. receives reward r_t
- Goal is to find the policy, π^* , that maximizes the agent's reward:

$$J|_{\pi} = E \left\{ \sum_{t=0}^{\infty} \gamma^t r_t \right\}$$

- Popular algorithms include: Q Learning, TD Learning, and Actor-Critic
- Policy gradient and natural actor critic are most popular in robotics



Simulation: Linear Quadratic Regulator

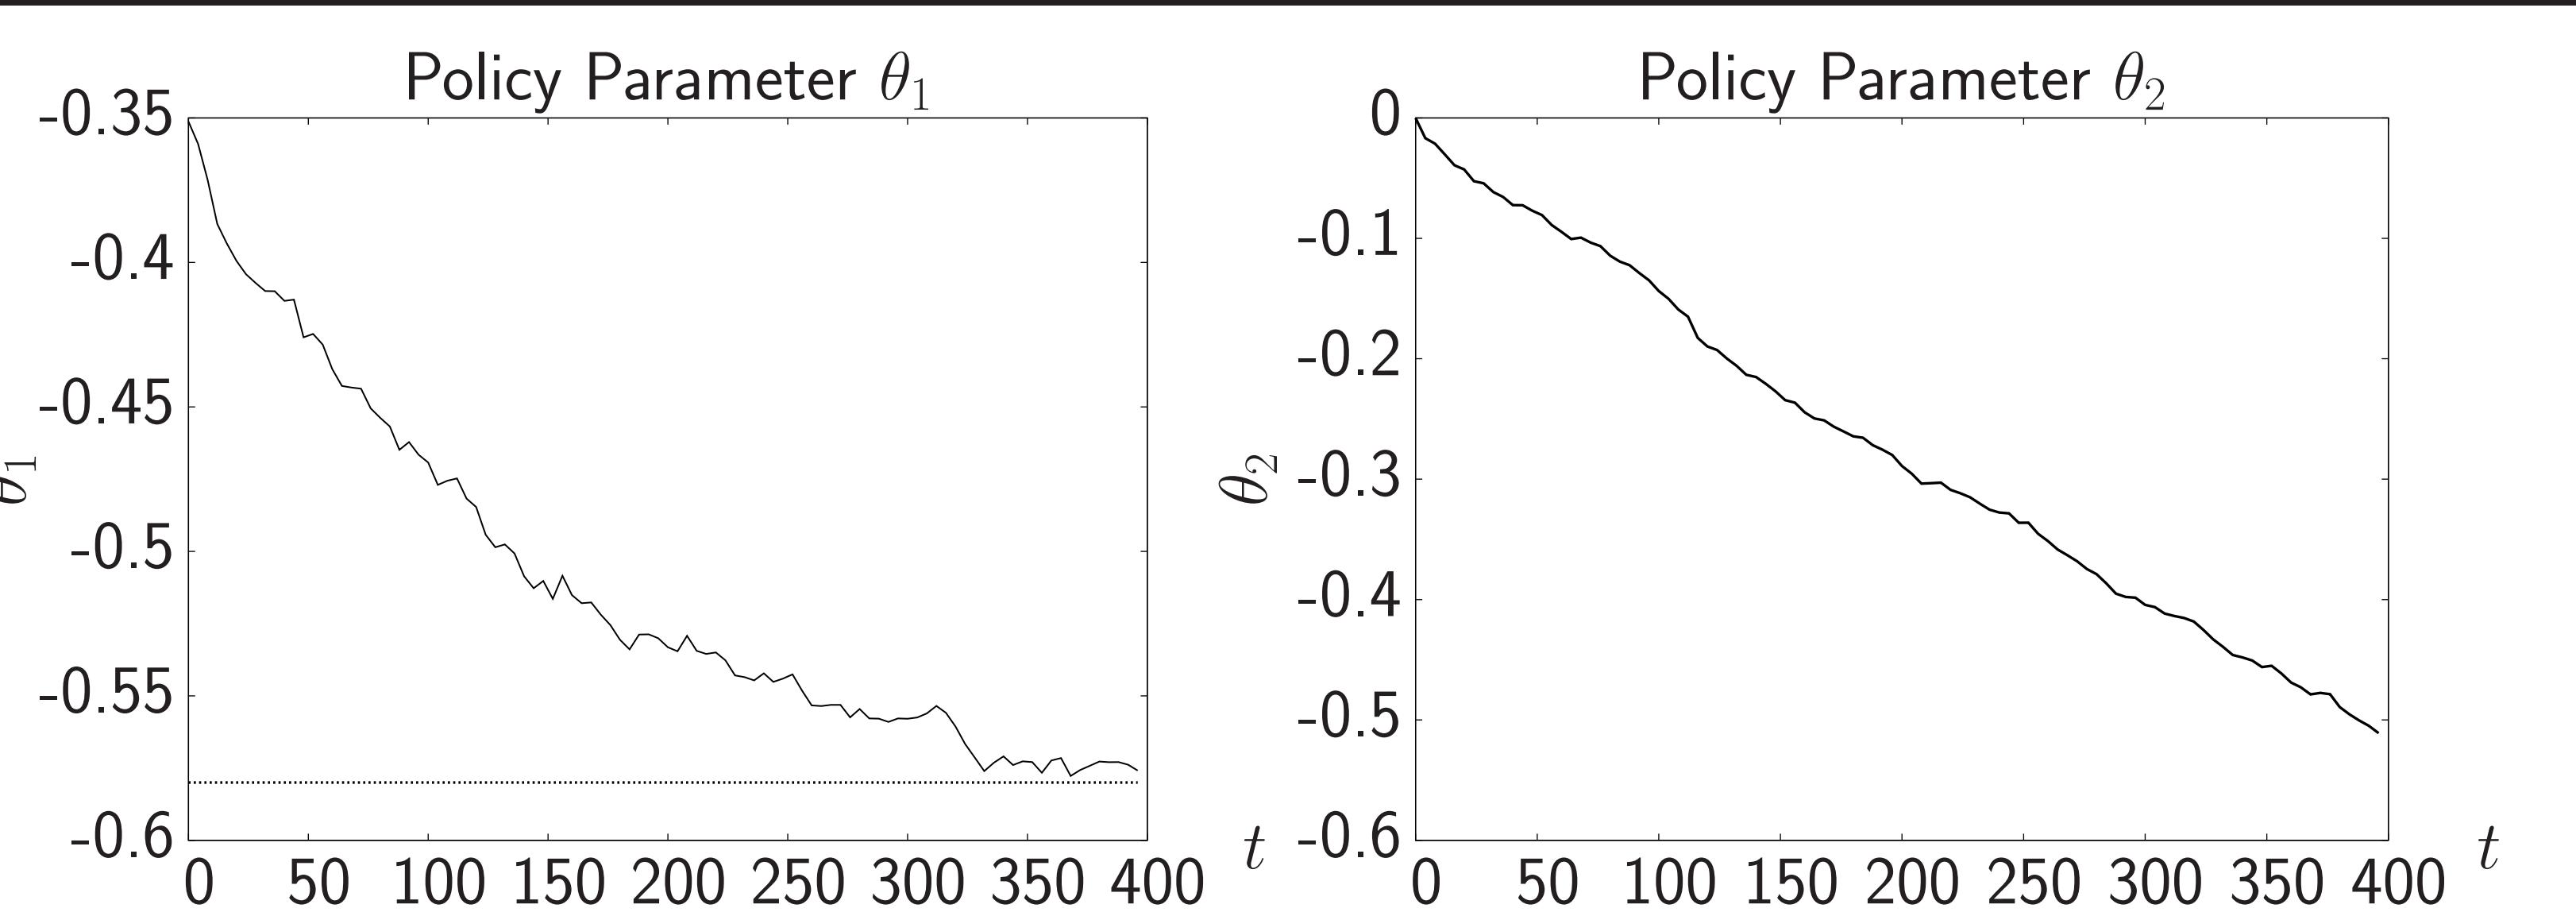


- Simple one-dimensional control theory problem with $s \in \mathbb{R}$, $a \in \mathbb{R}$, and the environment: [2]

$$\begin{aligned} s_{t+1} &= s_t + a_t + \text{noise} \\ R_t &= -s_t^2 - a_t^2 \end{aligned}$$

- Stochastic controller draws actions from $a_t \sim N(\mu_\theta, \sigma_\theta)$ where $\mu_\theta = \theta_1 s_t$ and $\sigma_\theta = [1 + e^{-\theta_2}]^{-1}$ [2]
- Vanilla policy gradient was simulated using central-difference estimator to compute gradients [3]

Results: Linear Quadratic Regulator Simulation



- Control theory proves that the theoretical optimum is at $\theta_1^* = -0.58$ and $\theta_2^* = -\infty$ [2]
- Vanilla gradient ascent quickly converges to $\theta_1 = \theta_1^*$ and continues moving towards $\theta_2 = \theta_2^*$

Policy Gradient

- Assumes the policy $\pi(s, a | \theta)$ is uniquely parameterized by θ
- Learning the optimal policy π^* is now simplified to learning the optimal θ^*
- For any given θ , $\nabla J|_{\theta}$ points in the direction of maximum increase
- Step the parameters by a small amount in the direction of $\nabla J|_{\theta_i}$:
$$\theta_{i+1} = \theta_i + \alpha \nabla J|_{\theta_i}$$
- Gradually decrease α until the algorithm converges to a local maximum
- **Problem:** This requires an empirical estimate of the gradient $\nabla J|_{\theta}$

Finite Difference Gradient Estimate

- Simple form of gradient approximation that requires no knowledge of π
- Slightly perturb θ with several small changes (i.e. $\Delta\theta_1, \Delta\theta_2, \dots$)
- Approximate $J(\theta + \Delta\theta_i)$ for each $\Delta\theta_i$ by averaging multiple roll-outs [3]
- Construct $\Delta\Theta$ and ΔJ such that $\Delta\theta_i = \theta_i$, $\Delta J_i = J(\theta + \Delta\theta_i)$, and

$$\nabla J|_{\theta} = (\Delta\Theta \Delta\Theta^T)^{-1} \Delta\Theta^T \Delta J$$

Likelihood Ratio Gradient Estimate

- Requires specific knowledge of the policy to estimate $\nabla\pi(s, a | \theta)$ [2]
- Estimates $\nabla J|_{\theta}$ without requiring the policy to be perturbed:

$$\nabla J|_{\theta} = E \left\{ \sum_{t=0}^{\infty} (r_t - b) \nabla \log \pi(s_t, a_t | \theta) \right\}$$

- Faster convergence than finite difference methods in noisy environments

Natural Policy Gradient

- Gradient is defined in terms of Euclidean distance:
$$\lim_{\Delta\theta \rightarrow 0} \left(\frac{\|J(\theta + \Delta\theta) - J(\theta) + \nabla J(\theta) \cdot \Delta\theta\|_2}{\|\Delta\theta\|_2} \right) = 0$$
- Parameters have arbitrary units, making the space highly non-Euclidean
- Vanilla gradient no longer points in the direction of maximum increase [3]
- Natural gradient is rotated to face the direction of maximum increase: [3]
$$\tilde{\nabla} J|_{\theta} = G_{\theta}^{-1} \nabla J|_{\theta}$$
- Dramatically better convergence rates than vanilla policy gradient [1]

References

- S.I. Amari.
Natural gradient works efficiently in learning.
Neural computation, 10(2):251–276, 1998.
- H. Kimura and S. Kobayashi.
Reinforcement learning for continuous action using stochastic gradient ascent.
Intelligent Autonomous Systems (IAS-5), pages 288–295, 1998.
- J. Peters and S. Schaal.
Policy gradient methods for robotics.
In *Intelligent Robots and Systems, 2006 IEEE/RSJ International Conference on*, pages 2219–2225. IEEE, 2006.
- Stephane Piskorski.
AR.Drone Developers Guide SDK 1.5.
Parrot, 2010.